

From Physical to Virtual Sensors (PVS)

Camilla Stormoen

INF-3983 Capstone Project in Computer Science ... December 2017



Abstract

W3 Whats wrong with the world? / motivation 1-3 setninger

Architecture - 1-3 setninger

Design- 1-3 setninger

Implementation - 1-3 setninger

Experiments - 1-3 setninger

Results - 1-3 setninger

Lessons learned/main conclusion - 1-3 setninger

Kutt heller etterpaa

This dissertation present/describe ...

Contents

Abstract	i
List of Figures	v
1 Introduction	1
1.1 Motivation	1
1.2 Contributions	2
1.3 Assumptions	2
1.4 Limitations	2
2 Related Work	3
2.1 Virtual Sensors	3
3 The Dataset	5
3.1 Dataset Directory Structure	5
3.2 Image Annotations	6
4 Architecture	9
4.1 Physical Sensors	9
4.2 Raw Data, Analytics and Labeled Data	9
4.3 Fusing Of Data	11
4.4 Fused Data	11
4.5 Virtual Sensors	11
4.6 User Interface	11
5 Design	13
5.1 Data Storage	13
5.2 Sensor Fusion	13
5.3 Virtual Sensors	15
5.3.1 Virtual Sensors	15
5.3.2 User Application	15
6 Implementation	17
6.1 Overview	17

6.2 Data Storage	18
6.3 Sensor Fusion	19
6.3.1 Raw Data Information Extraction	19
6.3.2 Excel files Extraction	19
6.3.3 Comparing	20
6.4 Virtual Sensors	20
6.4.1 Virtual Sensors	20
6.4.2 User application	21
7 Evaluation	23
7.1 Experimental Setup	24
7.2 Experimental Design	24
7.3 Results	24
8 Discussion	25
8.1 Architecture	25
8.1.1 Other solutions	25
8.2 Design	25
8.2.1 Other solutions	25
9 Contributions	27
10 Conclusion	29
10.1 Contributions?	29
11 Future Work	31
12 Appendix?	33
Bibliography	35

List of Figures

2.1	Figure illustrating the virtual rainfall sensor.	4
3.1	Images from the COAT dataset showing how a picture can contain colors or greyscale.	6
3.2	Figure showing a cropped screenshot of the structure of the directories.	7
4.1	Figure showing the system architecture.	10
5.1	Figure showing the system design.	14
6.1	Shows inconsistency between cells in rows in excel file.	19

/ 1

Introduction

- Mention focus on camera-sensors/data, and not other sensors?! *John M.: Hvorfor fokus på bilder - forenkle, begrense, starte et sted*
- Talk a little bit about COAT in general?

This project will develop an abstraction for virtual sensors, and do a prototype of the abstraction on a set of computers with physical sensors.

The purpose is to provide for a more powerful and flexible sensor in the COAT monitoring of the arctic tundra. As such, a fox feeding station is the usage domain to be used for the prototype.

1.1 Motivation

The motivation!

- W3
- Problem definition: This project investigated ... x, with the purpose of y.

The motivation behind this project is that no single sensor may cover the sensing needs, and that sensing needs can change rapidly over time. Consequently,

there is a need for sensor fusion, and allow for combining sensors at different computers.

1.2 Contributions

What was the contribution?

1.3 Assumptions

AVGRENSE, VIKTIG!! Something about motivation and stuff

1.4 Limitations

AVGRENSE, VIKTIG!!

- Mention focus on camera-sensors/data, and not other sensors?!

/2

Related Work

2.1 Virtual Sensors

A virtual sensor is a constructed sensor in contrary to a physical sensor [1]. They are used in place of the real sensors where they read real physical sensor data and calculate the outputs by using some processing models.

Previous research have focused on simple in-network data aggregation techniques and the sensor networks are often represented as a database. Two examples of such approaches are TinyDB [6] and Cougar [7], which enable applications to have a central point (base station) and create routing trees to funnel replies back to this root. The main focus on these approaches are operating intelligent in-network aggregation and routing to reduce the overall energy cost while still keep the semantic value of data high. In both examples, the data aggregation is specified using an SQL-like language. However, queries cannot be used to merge different data types, only homogeneous data aggregation is achievable. Their virtual sensors approach is offering a simple interface and heterogeneous in-network data aggregation. Yet, our virtual sensors have no SQL-like language or a database to store the aggregated data. Our work also relies on physical sensor data which is already in a data storage, and not directly from the physical sensors.

A virtual node[2] is a set of physical sensors that a programmer can interact with as a single sensor. Their virtual nodes are implemented using TinyOS [3] and their physical nodes are abstracted and specified using logical neighborhoods

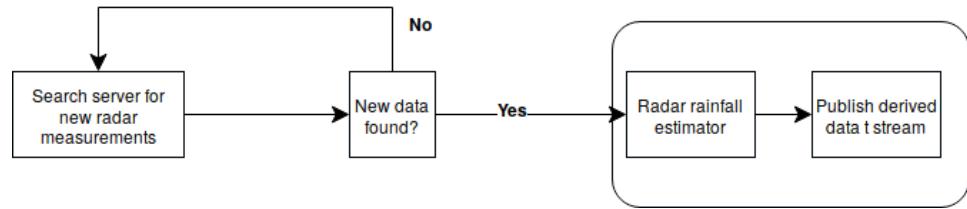


Figure 2.1: Figure illustrating the virtual rainfall sensor.

[4][5]. The nodes are combined in a logical neighborhood that are specified based on their characteristics by the programmer.

It is a prototype of a virtual sensor system for environmental observation with real-time customization of physical sensor data. The system can provide point-averaged radar-rainfall products at either temporal resolution of the radar or as a temporal average of a fixed time period, which is illustrated in Figure 2.1. In contrast, our virtual sensors system is running on-demand, not in real-time.
OGSÅ ANDRE ULIKHETER

/3

The Dataset

HVORDAN DE ER SAMLET, HVORDAN DE ER ORGANISERT?

The dataset is provided by COAT and contains over 1.6 millions of pictures taken from 2011 to 2015 by their camera traps stationed in the Arctic Tundra in Finnmark, Norway. The camera traps are placed at six different areas in the Arctic Tundra: Nordkynn, Ifjord, Komag, Nyborg, Stjernevann and Gaissene. The dataset contains pictures during night- and daytime. The cameras have infrared flash so the cameras can take pictures at nighttime. These pictures are without color while the pictures taken at daytime are with color, as shown in Figure 3.1.

Wildlife camera Reconyx

3.1 Dataset Directory Structure

The pictures in the dataset is structured in directories and folders. The dataset is first divided into different years going from 2011 to 2015. Each of these folders are then again divided into the areas of the 5 camera traps in the Arctic Tundra. Inside each of these folders specified by the area, they contain folders from each site inside the specific area. Figure 3.2 shows a cropped screenshot of the directories.



(a) Picture of ravens at daytime. (b) Picture of an arctic fox at nighttime.

Figure 3.1: Images from the COAT dataset showing how a picture can contain colors or greyscale.

(In total 1976 directories, men alle er ikke i bruk/relevant)

3.2 Image Annotations

Biologist working on the COAT project have detected animals in each image and made annotations stored in excel-files containing metadata and animal classification.

COAT provided annotations stored in excel-files with information about all of the images in the dataset that biologist working at the COAT project have detected. The biologist have used these annotations to see what animal was present at a specific time and to see statistics. The biologist have not made these annotations to go back in time and look at the images, and the dataset is then not made for...

The excel-files are structured in their own folder and each file is named as a combination of the camera traps area, it's year and the site. An example is "fotoboks2011_nordkynn_nordkynn.2011.xlsx".

Ta med hvordan/hva slags metadata som ligger i excel-filen? Evt i bildene.

The dataset contains the images metadata and animal classification, but did not contain any filenames or a filepath. We then had to find a method to find out which images correspond to the metadata in the excel-files. Fortunately, each image had metadata stored in Exchangeable Image File Format (EXIF). We recursively traverse the directories where images are located and read

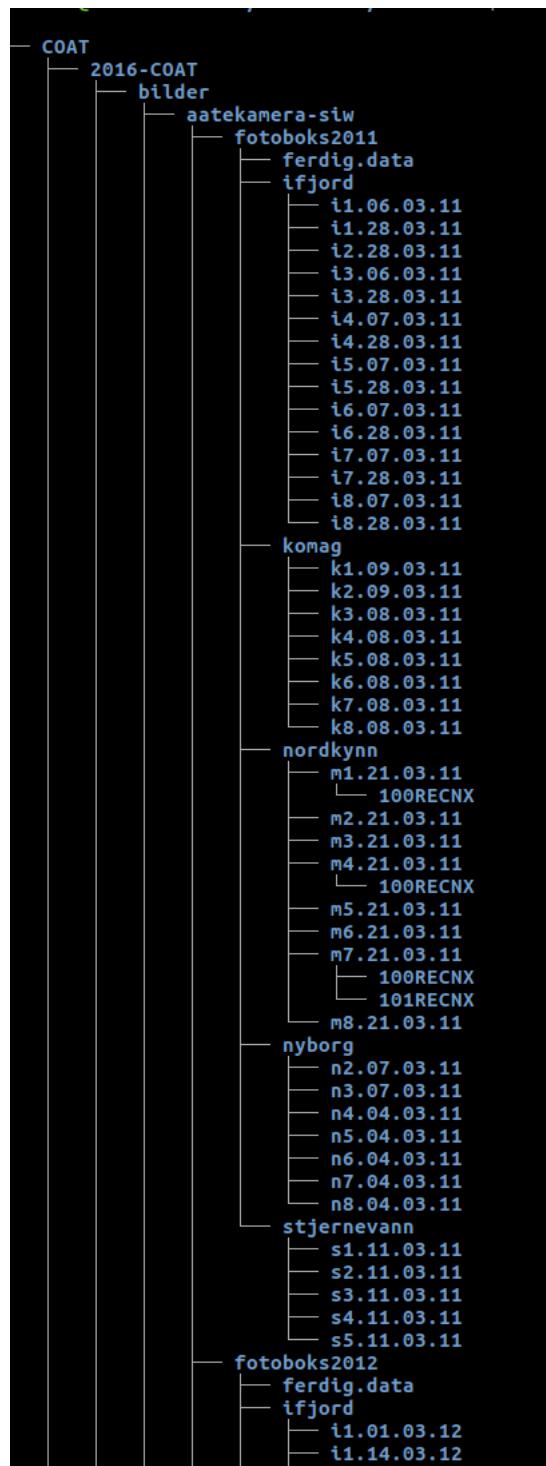


Figure 3.2: Figure showing a cropped screenshot of the structure of the directories.

date-time from the metadata and store it in a new CSV-file (Comma-separated Values). This was quite a time-consuming task because of the number of images to process.

Forklare mer hvordan dette er gjort i design/implementasjon. Skal jeg si noe som ".. explained in details in Chapter/Section ...??"

/ 4

Architecture

This chapter describes the architecture of the system. The components in the system are: physical sensors, raw data, analytics and classified data, fusion of data, fused data, virtual sensors and the user interface. The architecture of the system is presented in Figure 4.1. The arrows indicates the dataflow between each component in the system.

4.1 Physical Sensors

The physical sensors produce raw data. The raw data consists of images from different bait-camera sensors. The data also contains metadata about each image, as described in Chapter 3.

Hva mer??

4.2 Raw Data, Analytics and Labeled Data

The raw data contains images from the physical sensors and annotations in excel-files containing metadata about each image alongside labeled data, as described in Chapter 3.

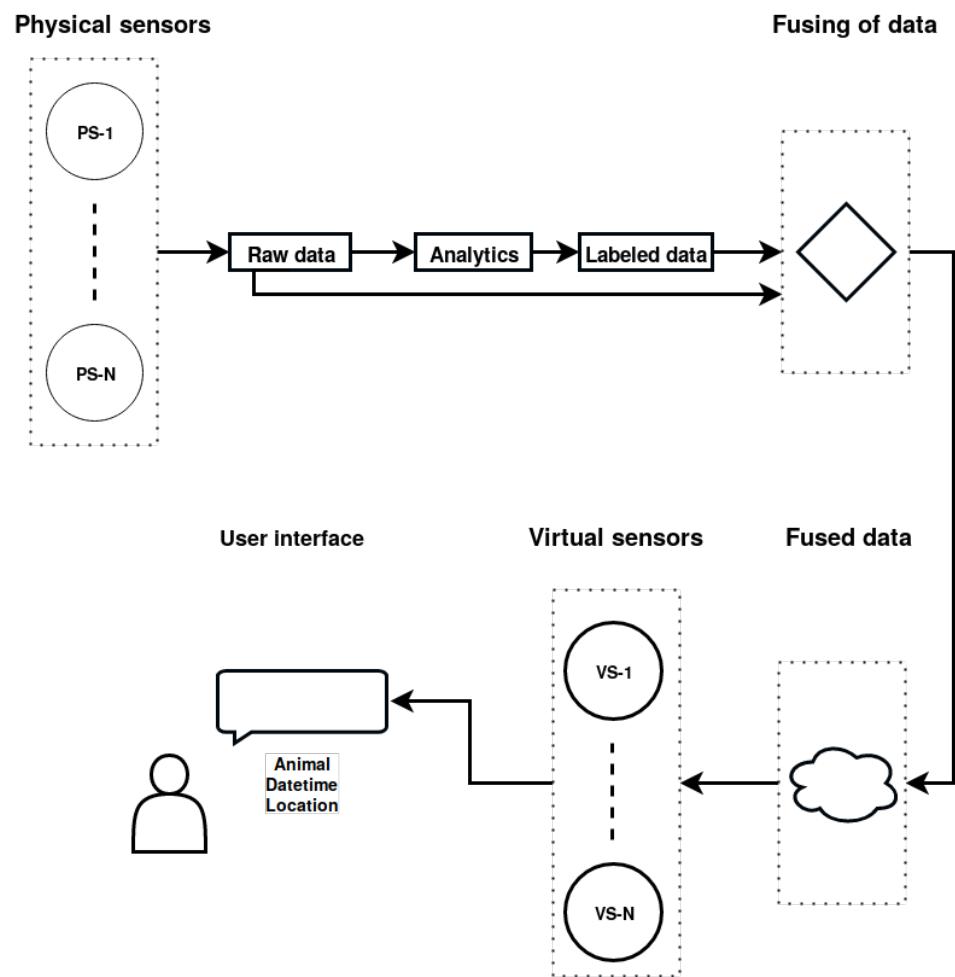


Figure 4.1: Figure showing the system architecture.

Hva mer??

4.3 Fusing Of Data

Get analytics and classified data from raw data (excel) and the raw-data and fuse/aggregate correlated/corresponding(?) data.

Hva mer??

4.4 Fused Data

The fused data retrieves it's data and on demand. The fused data is the physical sensors location combined with necessary metadata such as date-time, location, site, year, what kind of animal there as in the image and also how many animals there was/**and animal classification**.

4.5 Virtual Sensors

The virtual sensors are divided into multiple sensors where each virtual sensor represents a different animal that scientists are interested in, like ravens, red foxes, golden eagles or polar foxes.

UTDYP DE VIRTUELLE SENSORENE These virtual sensors have the same functionalities...?

The virtual sensors offers interaction between the fused data and the user interface. A user tells the user interface, described in Section 4.6, a task and forwards the task to the desired virtual sensor. The virtual sensor will then retrieve its results from the fused data.

4.6 User Interface

A user wants to retrieve fused data about animals. The user interacts with the virtual sensor user interface. This takes care of the interaction with the virtual sensor. When a user gives a command, the desired virtual sensor retrieve data from the fused data as described in Section 4.5, and presents the result back

to the user.

utdtype query språk - utdtype mer i design/implementation?

/ 5

Design

In this chapter we will look at the design of the system and present the design of each component of the architecture. Figure 5.1 shows the design of the system. The arrows indicates the communication lines. The dotted arrows show how the data storage is structured with files and pictures. **Få med direction of arrows/active side?**

5.1 Data Storage

The data storage contains raw data from the physical camera trap sensors and excel-files containing annotations for each raw data, as described in Chapter 3.

5.2 Sensor Fusion

As mentioned earlier in Chapter 3, the dataset did not contain any filename or a filepath. It only contained image metadata such as a pictures location, date-time, site, year and animal classification. We then had to find a method to find out which images correspond to which annotation in the excel-sheet.

We recursively traverse the directories where pictures are located and read

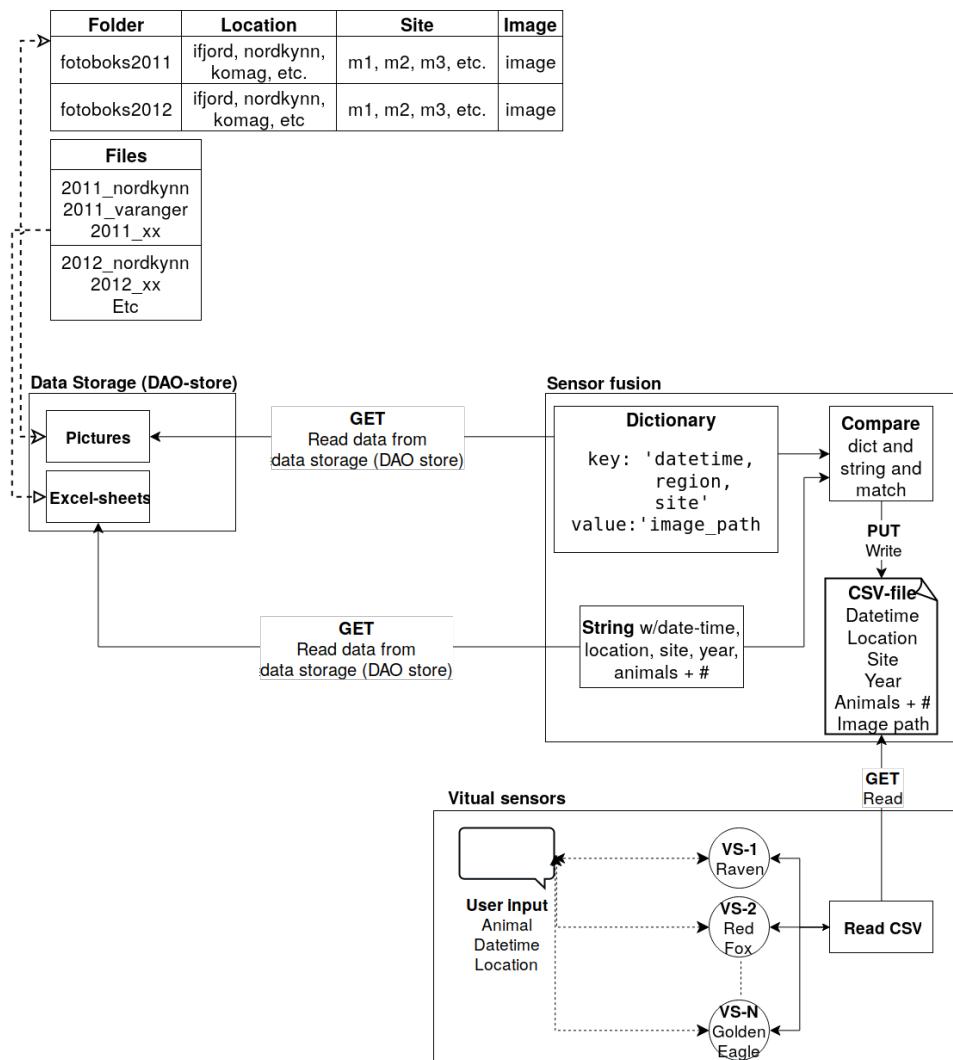


Figure 5.1: Figure showing the system design.

date and time from the metadata for each picture along with the annotation in the excel-sheet. The images and annotations that matched, is stored in a Python dictionary, like a key-value store, with date-time as key and image-path as value. As mentioned earlier, was this quite a time-consuming task because of the numbers of images to process and compare.

5.3 Virtual Sensors

5.3.1 Virtual Sensors

Divided into multiple virtual sensors depending on animal.

The respective virtual sensor must handle events from the user. Assume a user wants to find pictures of a red fox at a specific time. The user interface makes the respective virtual sensor aware of the task. The virtual sensor will then go to a list of all red foxes that are located from the fused data and search for the specific request. If the virtual sensor find pictures that are similar to the request from the user, one and one picture is visualized on the computer screen to the user.

The virtual sensors are divided into multiple sensors depending on animals...

The virtual sensors in this approach are functions in a sequential **system/program**. They could have been threads in a parallel program or the best solution would be if the virtual sensors where independent subprocesses because that would not affected other virtual sensors when a virtual sensor is added or removed from the system, as well as if one virtual sensor crashes.

Mer i implementasjon?

5.3.2 User Application

The user-application is the connection between the user and the virtual sensors. It gets input from the user and sends a request to the virtual sensor based on its input. The user-applications goal is to make an interface that is easy to use and to show images that the user want to see. The main menu is shown in Listing 5.1. The user only need to specify the following 3 parameters for the virtual sensor:

- **Animal:** The user can choose to see 3 different animals; raven, red fox and the golden eagle. The user can also chose to see pictures with no

animals in it. The reason that the user only can chose between these 4 categories is that these are the only virtual sensors that are implemented in this version of the system.

- **Date-time:** If the user wants a specific date-time the picture(s) was taken, he/she can specify the date and time in this section. If the user leave this field blank, all dates will included when the system search for pictures.
- **Location:** At last, the user can chose if he/she wants the picture to be located anywhere specific. The only valid arguments in this field is Nordkynn, Komag, Nyborg or Stjernevann since these are the only names in the CSV-file in the fused data.

Listing 5.1: Main menu

```
***** WELCOME! *****
*****
```

Write what you want to see
(Raven, RedFox, GoldenEagle):

Write when you want to see (blank **is all** dates):
E.g: 2011:03:24 **or** 2011:04:24 10:10:00

Where do you want the animal to be located
(nordkynn, varanger: komag, nyborg, stjernevann):

/6

Implementation

6.1 Overview

Threads, data structures, language...

This chapter will elaborate on how we implemented the system, general implementation requirements, issues and choices.

The system is implemented and written in the high-level programming language Python 2.7¹. This was a practical ... reason 1, reason 2, reason 3... We made this choice because Python is object oriented and it offers different frameworks that are relevant for this implementation.

Drawback with Python - not that good at parallel/concurrency..

In the rest of this chapter we will first look at the data storage in Section 6.2, then we will take a look at the sensor fusion in Section 6.3, and at last we will describe the virtual sensors in Section 6.4.

libraries used in this implementation:

1. <https://www.python.org/>

Exchangeable Image File Format - EXIF

Exchangeable Image File Format (EXIF) is a standard that specifies the formats for images, sound, scanners and other systems recorded by digital cameras. This standard consists of image information of the EXIF image file or the EXIF audio file.

Python have a module called EXIF 2.1.2², to extract the metadata(image information from the jpg-files.

Pandas Dataframe

Pandas³ is an open source Python Data Analysis Library which provides high-performance and easy-to-use data structures. It is a NUMFocus⁴ sponsored project which will help ensure success of the development of pandas as a open-source project in world class. A dataframe⁵ is a two-dimensional, dynamic tabular data structure with labeled axes (rows and columns), just like an excel-file.

OpenCV

OpenCV V2.4.9.1⁶ is an open-source library which provides the functionality for loading and preprocessing images.

6.2 Data Storage

The data storage contains folders with labeled and sorted animals, but these images consists of no metadata so we don't have any information of when the image is taken. Therefore, the solution is to traverse the whole data storage and find the original images with metadata and compare these to the excel-files containing information like date, time, location and site. **HOW DID I TRAVERSE**

2. <https://pypi.python.org/pypi/ExifRead>
3. <http://pandas.pydata.org/>
4. <https://www.numfocus.org/open-source-projects/>
5. <https://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.html>
6. <https://opencv-python-tutroals.readthedocs.io/en/latest/>

6.3 Sensor Fusion

HVORDAN EXTRACTED JEG DATA FRA EN DATAFRAME?!

6.3.1 Raw Data Information Extraction

The raw data in the data storage is traversed by walking the OS directories. If the traversed file ends with JPG, the image information in the EXIF is being extracted and the datetime value is read. This datetime along with the images site, region and image path is stored in a Python dictionary, where the datetime, site and region is key and the image path is the value.

6.3.2 Excel files Extraction

The excel files contains annotations for each image. To read excel files, we use Pandas Dataframe, described in Section 6.1. From the excel files, we extract information like date, time, animals, region, site and year. The necessary information about each animal is stored in a string.

During implementation we came across a problem with inconsistency between cells in the same row. The problem was that the datetime cells information was inconsistent with the information in cells day, month, year, hours and minutes, shown in Figure 6.1. Our solution was to alter and concatenate the correct values to make a new datetime.

Another problem was that the cells containing column-names was different. One example is that month was written "moth" in some excel files and "month" in others.

Solution: overwrite moth to month? Not done, though about it while writing this now... Only tried with two excel files that contains "moth", time restrictions and wanted to find some pictures instead of nothing..Excel-file har også æøå som ikke vises i Ubuntu sin excel libre offie calc.. Kanskje

23	3	2011	7	40	0m7.11	T	3/22/2011
23	3	2011	7	50	0m7.11	T	3/22/2011
23	3	2011	8	0	0m7.11	T	3/22/2011
23	3	2011	8	10	0m7.11	T	3/22/2011
23	3	2011	8	20	0m7.11	T	3/22/2011
23	3	2011	8	30	0m7.11	T	3/22/2011
23	3	2011	8	40	0m7.11	T	3/22/2011
23	3	2011	8	50	0m7.11	T	3/22/2011

Figure 6.1: Shows inconsistency between cells in rows in excel file.

bare hos meg, men burde gått igjennom og hatt en consistency mellom navn.. og unngå æøå. Fikse på raw data

A different problem was that some column-header names was written with a big letter like "Time" in some excel files and other contained "time". Our solution was to make every column header into lowercases using Pandas Dataframe own function, `pandas.Series.str.lower()`.

6.3.3 Comparing

Disse sammenlignes og de som matcher blir skrevet til en ny CSV fil med nødvendig info for å finne bilde.

6.4 Virtual Sensors

6.4.1 Virtual Sensors

ikke threads ettersom man evt vil addere flere sensorer og unnga a starte alle sensorer på nytt igjen.. Er de virtuelle sensorene servere eller client/publisher?

The virtual sensors are implemented as functions, not threads or independent subprocesses, as mentioned in Section 5.3.1. The reason for this choice is because ...

Valgte å fokusere på 3 dyr (og alle dyr/ingen dyr i bilde) ettersom de var de som det var flest av i excel-arkene.. and time constraints.. ville ha noe å vise enn ingenting

Write about that you don't have all animals, only the ones who is most of in csv-file, not every animal is presented in the dataset..

The OpenCV library is used to visualize/show pictures to the user in the user application. The library read the pictures it is given and display these on the screen to the user.

John M.: lesere vil kanskje tenke at dette er overkill. Kanskje bare gjøre et poeng at dette bare ble valgt på grunn av kjennskap til det / bruk i andre prosjekter og at nøyaktig hvilket verktøy som brukes til visualisering ikke var fokus.

6.4.2 User application

The user application is implemented as a interface in the terminal... There is currently no logging of a user's action. Nor is it possible for more than one user to be connected (in the same application) at a time. Because of lack of time, we have not thought about implementing this...(One user - one application, another user run another new application). Restricted to 3 animals and 4 places..



7

Evaluation

This chapter describes the experimental setup...

metrics, define (CPU, memory, latency.), benchmarks (micro, kernel... How to measure, where done, PSEUDOCODE

NOTATER:

- Time: Finding folders and metadata takes: 1:43:13.488799 (103,2167 minutes), Reading excel file takes: 0:00:17.413845 (0.2833 minutes), Comparing takes: 4:43:30.705587 (283.5 minutes), Overall time is 6:27:01.608355 (387.0167 minutes). Med alle bilder m/metadata og hele fotoboks2011_nordkynn_nordkynn.2011.xlsx.
- New time: Finding folders and metadata takes: 1:46:17.406581 (= 106.2833 minutes) Reading excel file takes: 0:01:07.686779 (= 1.1167 minutes) Comparing takes: 19:03:11.177869 (= 1143.1833 minutes) Overall time is 20:50:36.271415 (= 1250.6 minutes) Med alle bilder m/mETADATA og hele nordkynn og varanger
- "Concurrent" python find_folder_ok_concurrent.py Finding folders and metadata takes: 1:30:50.250732 (= 90.8333 minutes) Reading excel file takes: 0:00:18.127597 (= 0.3 minutes) Comparing takes: 4:18:33.368199 (= 258.55 minutes) Overall time is 5:49:41.746759 (= 349.6833 minutes)
– Pool(processes=4) FRA fotoboks2011_nordkynn_nordkynn.2011.xlsx

- "Concurrent" igjen Finding folders and metadata takes: 1:44:18.788013 (= 104.3 minutes) Reading excel file takes: 0:00:17.269817 (= 0.2833 minutes) Comparing takes: 3:54:23.483186 (= 234.3833 minutes) Overall time is 5:38:59.541248 (= 338.9833 minutes) – Pool(processes=100)
FRA fotoboks2011_nordkynn_nordkynn.2011.xlsx

This chapter describes the experimental setup and metrics used to evaluate the implemented system.

7.1 Experimental Setup

All experiments were done on a Lenovo ThinkCenter with an Intel® Core™ i5-6400T CPU @ 2.20GHz × 4, Intel® HD Graphics 530 (Skylake GT2), 15,6 GiB memory and 503 GB disk. It ran on Ubuntu 17.04 64-bit.

7.2 Experimental Design

7.3 Results

In this section we will discuss the results of the testing described in the sections above...

What does the result say? Each experiment, result, meaning

/ 8

Discussion

8.1 Architecture

8.1.1 Other solutions

- Virtual sensors probably uavhengige prosesser, ikke threads ettersom man evt vil addere flere sensorer og unngå å starte alle sensorer på nytt igjen..
- Updates of "database" in background, not on demand -> future work?

8.2 Design

8.2.1 Other solutions

- hash-map instead of multiple lists with different animals. hash on animal and/or place
- parallel/concurrent program when doing metadata-collecting and reading from excel-files and writing to csv. Python - bad experience with threads/subprocess, golang probably better?

/9

Contributions

Combine with conclusion??



10

Conclusion

10.1 Contributions?



11

Future Work

/12

Appendix?

readme, source code, dataset measurement RAW

Bibliography

- [1] S. Kabadai and A. Pridgen and C. Julien, *Virtual Sensors: Abstracting Data from Physical Sensors*, 2006, in *2006 International Symposium on a World of Wireless, Mobile and Multimedia Networks(WoWMoM'06)*, 6 pp.-592.
- [2] Ciciriello, Pietro and Mottola, Luca and Picco, Gian Pietro, *Building Virtual Sensors and Actuators over Logical Neighborhoods*, 2006, in ACM, 19–24. <http://doi.acm.org/10.1145/1176866.1176870>.
- [3] Hill, Jason and Szewczyk, Robert and Woo, Alec and Hollar, Seth and Culler, David and Pister, Kristofer, *System Architecture Directions for Networked Sensors*, 2000, in *ASPLOS-IX: Proc. of the 9 nt Int. Conf. on Architectural Support for Programming Languages and Operating Systems*, pages 93–104.
- [4] Mottola, Luca and Picco, Gian Pietro, *Logical Neighborhoods: A Programming Abstraction for Wireless Sensor Networks*, 2006, in *Distributed Computing in Sensor Systems: Second IEEE International Conference, DCOSS 2006, San Francisco, CA, USA, June 18-20, 2006. Proceedings*, pages 150–168. https://doi.org/10.1007/11776178_10.
- [5] Mottola, Luca and Picco, Gian Pietro, *Programming Wireless Sensor Networks with Logical Neighborhoods*, 2006, in *Proceedings of the First International Conference on Integrated Internet Ad Hoc and Sensor Networks*, series = InterSense '06, pages 150–168. <http://doi.acm.org/10.1145/1142680.1142691>.
- [6] Madden, Samuel R. and Franklin, Michael J. and Hellerstein, Joseph M. and Hong, Wei, *TinyDB: An Acquisitional Query Processing System for Sensor Networks*, March 2005, in *ACM Trans. Database Syst. Vol. 30, Nr. 1*, pages 122–173. <http://doi.acm.org/10.1145/1061318.1061322>.
- [7] Yao, Yong and Gehrke, Johannes, *The Cougar Approach to In-network Query Processing in Sensor Networks*, September 2002, in *SIGMOD Rec.*, Vol. 31, Nr. 3, pages 9–18. <http://doi.acm.org/10.1145/601858.601861>.

- [8] David J. Hill and Yong Liu and Luigi Marini and Rob Kooper and Alejandro Rodriguez and Joe Futrelle and Barbara S. Minsker and James Myers and Terry McLaren, *A virtual sensor system for user-generated, real-time environmental data products*, 2011, in *Environmental Modelling & Software*, Vol. 26, Nr. 12, pages 1710 - 1724. <http://www.sciencedirect.com/science/article/pii/S1364815211001988>.